

echo.

Agent continuity across people, sessions,
and platforms. Like nothing else.

For teams and super individuals.

echoknows.com • kobe@iditor.com

Put every Meta employee on Echo today.

By year end, Meta saves **≈\$0.6–1.0B**

And the work gets better, and faster.

Estimate • the math comes later

\$4M SAFE Seed



Circle area = valuation, smallest to largest, counter-clockwise from the left ~ estimated from latest 2026 round (valuation undisclosed)

Agent coding exploded.

GitHub, where the world keeps its code, **broke** first.



Vlad Fedorov GITHUB CTO

"By February, it was clear that we needed to design for **30X today's scale**."

Dramatized • real quote: GitHub Blog 2026-04-28


We need 10x capacity...

No, no, no –

30x!



The agent-era surge agent-authored PRs • Sep 2025 = 1x



14_B

commits this year

~1B all last year • 275M/week now

257

incidents in one year

48 major outages • ~one a week

Not one hot product. The whole stack can't keep up with agents.

Agents are moving from coding to everything.



\$9B → \$65B

Anthropic run rate – in just 7 months



\$20B → \$40B

OpenAI run rate – doubled in a year



25M

Codex active users · 1M in February
+5M in the last 10 days 2026.08.31



\$207M

avg AI budget, large US firms – nearly 2×

ENTERPRISE CODEX WEEKLY ACTIVES BY FUNCTION · FEB 2026 = 1× · LOG SCALE



And that's only trend one.

Three obvious things.

\$105B

TREND 01 • OPENAI + ANTHROPIC

HAPPENING NOW

Companies already pay for the work agents do.

OpenAI **\$40B** + Anthropic **\$65B** – **\$105B** combined 2026 run rate, up from under \$30B a year ago. Budgets are shifting from buying software to buying machine labor – the explosion on page 03.

a16z 81%

TREND 02 • A16Z • JAN 2026

HAPPENING NOW

No model vendor holds the enterprise. Using several is the norm.

a16z surveyed 100 Global 2000 CIOs: **81%** run **3+ model families**, up from 68% less than a year ago.

👉 Ramp spend data: **79%** of businesses paying Anthropic also pay OpenAI.

0

TREND 03 • STILL EMPTY

STILL EMPTY

Continuity **doesn't exist yet.**

Zero platforms let context leave their own walls. OpenAI's memory works only inside OpenAI, Anthropic's only inside Anthropic – **no platform will build its own exit.**

Context is stuck behind vendor walls – **every session rebuys the same work.**

Your most important context is now agent sessions.




TODAY Everyone keeps rebuilding the same context.

MANAGED Everyone shares one context.

10x more efficient

Agent memory has three flaws. Complex work stalls.


01



Markdown never works.

Kept by hand, it never keeps up.

02



Breaks at every handoff.

Codex can't see what Claude knows.
Yesterday never reaches today.

03

99.5%

Brute-force context.

Every session rereads everything.



"Every 6 months delete your CLAUDE.md. Delete your skills. Delete your hooks."

Boris Cherny · Head of Claude Code, Anthropic · YC Startup School 2026



06:54

Meta is paying tuition: tokens don't buy team output.

 Meta 2026

SPEND

\$4-6B

to  Anthropic EST.

OUTPUT

→ **220%**

more code changes, YoY

COST

74% → **55%**

employee sentiment, same period

40%

more major incidents

70%

more time putting out fires

The bottleneck isn't intelligence.
It's work that doesn't connect.



57%

use agents
in many stages



16%

connect them
across teams

How Echo saves Meta a billion, starting now.



01 **Every session pours into Echo.**

Every tool, every teammate, every agent.

02 **Memory flows back out.**

Token spend becomes a second brain.

03 **\$1B+ /yr** saved at  Meta
\$40B ↗ industry-wide

What one agent just learned, every agent uses in real time.



01 Capture

Plugs into the agents teams already use.

02 Remember

Why, who, where from, and who may see it.

03 Resume

The next agent starts on turn one.


-60-70% total token spend
no repeated context

-50% time per task
fewer, faster turns

3.2x effective output
1.5 → 4.8 a day

2.4x projects in parallel
7 → 17, same team

In every team, Echo carries the handoff.




Agent use since February **x108** Legal **x41** Sales • Recruiting **x26** Marketing

Every new team is a new handoff.


Nobody has built this yet: too new, too granular, too hard.

01 Too new




Teams where everyone works through agents only appeared this year.

02 Too granular



Existing memory keeps facts about a person, not a team's work in motion.

03 Too hard



Production needs all five, at the same time.

That is why the gap is still open, and it is exactly what Echo is built to do.

Others each solve a piece; Echo holds all of it at once.

WHAT A TEAM NEEDS	 ChatGPT Business	 Claude Enterprise	 GitHub Copilot	 Glean	 Supernormal	 Mem0	echo.
 Team work in progress Captured from agent sessions, not facts about you	—	◐	—	◐	◐	◐	●
 One memory, every agent Claude Code · Codex · ChatGPT · Cursor	—	—	—	◐	●	●	●
 Encryption Even the provider can't read it	—	—	—	—	—	—	●
 Fine-grained privacy control The owner shares item by item	—	—	—	◐	◐	◐	●
 Auditable and traceable Who wrote it, who read it, where it came from	—	—	◐	◐	◐	◐	●


Moat

01 Five years early




We started the engine in 2021, a year before ChatGPT. Today it ranks #1.

02 Deeper with use



Every day of team work makes the company's memory deeper.

03 Even we can't read it



Only the owner holds the key, so companies trust Echo with what matters most.

The longer a team works with Echo, the harder Echo is to replace.


3 weeks of founder-led sales: 2 signed, 3 in talks.

2 Signed

paying · 25 people


3 In talks

788 seats in reach



No sales team yet. The \$4M builds one.

10,000 paid seats gets us to \$600M.



Unlimit is the first 650 seats. Then we repeat the motion.

28 companies are hiring to solve this: $\approx 85,000$ seats.



Baseten



Abridge



Legora



Commure



Vanta



voize



Justworks



Grafana



Ambience



Benepass



Scale AI



Alpaca



Ramp



Notion



Stripe



Canary



Instacart



Datadog



Taktile



Tailscale



Forter



Deepgram



Gusto



Anthropic



Abacum



Clara



Shepherd



Aera Technology

Context sharing is the wedge, not the whole market.

The \$4M has one job: 10,000 paid seats in 18 months.




10,000 seats × \$100 / month = the \$10M+ ARR on page 15.

The risk isn't the technology; it's the sales.

01 Tech: models can't replace it

- Companies pay agents for finished work.
- Real work is variable and conditional.
- It almost never finishes inside the model.
- It needs a memory outside it. We build that.

02 Sales: every deal needs on-site setup



The first deal costs the most; each one after gets cheaper.



Reusable harnesses
one per industry



The \$4M boosts sales
more engineers in the field


The tech holds. The \$4M boosts sales.

Agent Continuity for Enterprises.



The floor is already priced near \$10B. The layer above it is where we're going.

In the agent era, Echo is the medium of thought.



Six people, full-time, AI-native

Strong in mathematics, statistics, social products and systems engineering.



Kobe Wang

Founder & CEO · ex-Roblox Social, 100M+ DAU



Biwei (Betti) Tao

Microsoft AI · ex-Amazon Alexa · UC Berkeley Statistics



In the agent economy, a company's context is its most valuable asset. We are building its engine.

Raising a \$4M SAFE.



2 signed • 3 in talks

in 3 weeks, founder-led



Rolling close

join anytime



Boost sales

forward-deployed engineers

10,000 paid seats → \$10M+ ARR → \$600M

IN 18 MONTHS

OX \$60M
SageOx

 ~\$110M
Mastra

 ~\$130M
Mem0

 ~\$250M
Dust

 \$600M
Engram

 \$1.5B
Granola

 \$7.2B
Glean

 \$600M
Echo target

kobe@iditor.com