

THE INDUSTRY, AS PRICED
BY CAPITAL

\$4M SAFE

Build the sales pipeline, go to market

THE SECOND BRAIN FOR AGENT-NATIVE COMPANIES

1st

in Accuracy

95.8% LongMemEval

1st

in Token Efficiency

-63% input tokens

1st

in Deployment Speed

3 paying enterprise
customers

The agent era is just starting. Serve enterprises, stay alive, then catch the wave.



These valuations aren't buying revenue. They're buying a head start.

50 COMPANIES TRACKED · 89 ROUNDS · \$11.18B CITABLE POST-MONEY · ONLY 2 PUBLIC ARR FIGURES — THE MONEY ARRIVED FIRST, THE REVENUE HASN'T · PUBLIC DATA, AS OF 2026-08-26

WELLINGTON LED



\$7.2B

~\$100M ARR at raise · 72x

The **only company of the 50** with real revenue at pricing — and the cheapest way in **still cost 72x revenue**.

ROUND 2025-06 · ARR ~\$100M THEN, \$300M NOW

CLOSED 50% ABOVE THE LEAK




\$1.5B

\$0 public ARR

In talks at \$1B in January. Closed at \$1.5B two months later — **the price moved 50% in eight weeks**, with no public revenue.

FORBES 2026-01 · ROUND CLOSED 2026-03

SEQUOIA · GENERAL CATALYST · AMPLIFY



\$600M

13 people · \$0 revenue

Day one out of stealth: **\$98M at \$600M**.
Thirteen people, no revenue — and three top funds all in.
= the dashed \$600M ring on page 01

ROUND 2026-06 · ZERO DISCLOSED ARR

A pre-revenue seat here costs \$600M. Echo's costs \$30M today.

Agent usage is exploding. And it's spreading.

\$9B → \$65B

Anthropic annualized run rate – in seven months

\$40B

OpenAI run rate – doubled this year


20M

people now use Codex / ChatGPT Work agents

\$207M

avg US enterprise AI budget – doubled in a year • KPMG

ENTERPRISE AGENT USAGE BY FUNCTION • INDEXED, FEB 2026 = 1× • LOG SCALE



GROWTH IN ENTERPRISE AGENT USE BY BUSINESS FUNCTION, FEB → AUG 2026 • PUBLIC USAGE DATA

And that's only trend one.

Two trends are certain.

\$105B

TREND 01 · OPENAI + ANTHROPIC
ARR

HAPPENING NOW

Companies are paying for agents' work.

OpenAI \$40B + Anthropic \$65B – \$105B combined 2026 annualized revenue, up from under \$30B a year ago. Budgets are moving from buying software to buying machine labor – that's the explosion on page 03.

81%

TREND 02 · A16Z

HAPPENING NOW

No platform can hold users. Moving is the norm.

81% of the Global 2000 run 3+ model families · 79% of OpenAI's business customers also pay for Claude.

0

TREND 03 · STILL MISSING

STILL MISSING

Continuity **doesn't exist.**

Zero platforms let context leave their walls. OpenAI's memory works only in OpenAI, Anthropic's only in Anthropic – no platform will ever build its own exit.

All those valuations are bets on context. Theirs can't leave the building. **Ours can.**

People are paying their agents to forget.

BREAK 01 · TIME



Your own agent forgets.

Close a session – tomorrow you **explain everything again.**

99.5%

of agent tokens are **input** – re-explaining, not producing · our 828 production sessions, 6.8B tokens

BREAK 02 · TOOLS




Your agents don't talk to each other.

Claude has no idea what Codex just did. You carry context between them **by hand.**

66%

of everything agents read is a **re-read** – finding what they already found

BREAK 03 · TEAM



Your team repeats each other's work.

Their agent spends the morning figuring out **what yours already figured out.**

~\$40B/yr

what enterprises burn on this today – **up from ~\$11B in all of 2025**

(\$65B ANTHROPIC + \$40B OPENAI, JUL 2026) × $\frac{2}{3}$ ENTERPRISE × 60% RE-EXPLAINED INPUT · 2025: \$18B MODEL LAYER × 60% · ECHO ESTIMATE

Everyone re-explains. Everyone re-pays. **Every day.**

Echo fixes all three. It's live today.

FIX 01 • TIME



Next session starts where the last one stopped.

Warm start – no re-explaining.

–63%

input tokens – the re-explaining, gone

FIX 02 • TOOLS



All your agents share one memory.

Install and it works – no integration project.

–73%

repeated searching – agents stop re-finding

FIX 03 • TEAM



Your team stops repeating itself.

Figured out once – reused by everyone.
Permissions and audit built in.

5x → 1x

a 5-person team rebuilds the same context **five times** – **Echo builds it once**. The gap compounds with every seat

–63% / –73% MEASURED IN PRODUCTION, 828 SESSIONS • TEAM REUSE = MODEL, N REBUILDS → 1 • PUBLISHED AT ECHOKNOWS.COM/RESEARCH

The \$40B on page 05 is optional. This is the proof.

GitHub for company progress

Every result keeps its root, lineage, and people.



Tokens are a liability. Memory you can dispatch is the asset. Echo turns one into the other.

We're #1 because we started five years early.

LONGMEMEVAL · HARD SET · PUBLISHED SCORES

1st Echo

95.8%

2 Hindsight

94.6%

3 mem0

94.4%

The highest published score. **Three months** – nobody has passed it. And the engine that holds the lead keeps iterating.

WHY NOBODY PASSES IT



Born from a Zettelkasten knowledge engine.
Native – not bolted on.

LUHMANN: 90,000 CARDS → 70 BOOKS · HIS "SECOND MEMORY" – WE BUILT IT FOR AGENTS



NATIVE EXECUTION – ACCURATE AND FAST

Compression and recall built **directly on the Zettelkasten engine** – scalable, auditable, manipulable **by design**, not by patch.



DEEP IN THE SCENARIO

Five years **inside real agent workflows**, not behind an API. That's why the 828-session dataset exists **only here**.



STRAIGHT-LINE – NO HUMAN IN THE LOOP

Smooth, simple, **fully automatic** – no tagging, no curation, no integration. **Memory that feeds itself**.

PUBLISHED 2026-05 · ECHOKNOWS.COM/RESEARCH

2021 · ZETTELKASTEN ENGINE

→ AGENT MEMORY

5 YEARS · VERIFIED · ITERATED

2026 · #1 · 828 SESSIONS LIVE

A score can be chased. An architecture can't be retrofitted.

The most important context in the company now lives inside its agents.

Echo is the control plane for what agents know, share, and act on.

HARD PASS CRITERIA	 ChatGPT Business	 Claude Enterprise	 Glean	 Engram	 Supermemory	 Mem0	echo.
Live agent work state GOAL · DECISIONS · CONSTRAINTS · OPEN LOOPS · NEXT MOVE	—	—	—	—	—	—	●
Portable user-owned memory USER ACCOUNT · EXPORT/DELETE · SURVIVES VENDOR SWITCH	—	—	—	—	◐	◐	●
Same memory across agents CHATGPT · CLAUDE · CODEX USE ONE MEMORY	—	—	—	◐	●	◐	●
Ready without integration NO API · NO RAG STACK · NO MANUAL DOCUMENTS	●	●	◐	—	◐	—	●
Provider cannot decrypt USER-HELD KEY · ZERO-KNOWLEDGE BY DESIGN	—	—	—	—	—	—	●
Source-to-decision trace SOURCE · OWNER · REASON · TIMESTAMP · CORRECTION	◐	◐	◐	—	◐	◐	●
Memory-level team controls PRIVATE → TEAM CONSENT · SCOPE · AUDIT · REVOKE	—	◐	◐	—	◐	◐	●

Only Echo passes every test—without a custom stack. This table is already the past – the next test is making agents finish the work.

2 signed. 4 in conversation. Three weeks.

Founder-led, before a sales team. Statuses shown exactly as they stand today.

SIGNED

2 companies

PAID DESIGN PARTNERS



ON1 Force
15 EMPLOYEES

• WEEK 0



Executive
Intelligence
10 EMPLOYEES

IN CONVERSATION

4 companies


ACTIVE FOUNDER-LED PIPELINE




etalytics
74 EMPLOYEES



Unlimit
650 SEATS



Xola
64 EMPLOYEES



Sound of
Molecules
3 EMPLOYEES

• WEEK 3

Two contracts signed before building a sales team.

Four more are already moving: industrial AI, fintech, travel software, and a biology / new-drug research team.

The motion exists. The \$4M scales it.

The \$4M goes to sales. That's it.

We hire forward-deployed engineers to reach companies **already hiring people to deploy and manage agents.**

RECENT HIRING SIGNALS • FDE / AGENT WORKFLOW

Public job posts • Jul-Aug 2026



 OpenAI	 Anthropic	 Salesforce
 Ramp	 Commure	 Gecko Robotics
 Lindy	 Runway	 Unlimit

120 target accounts → 15 paid audits → 8 pilots → 4-6 annual contracts

\$600M is the target. It's not the ceiling.

10,000 paid seats gets us to \$600M.

Unlimit's CCO publicly called on the CTO to start working with Echo. Then we repeat the motion.



Unlimit is the first 650 seats. The 28-company wall is only the first wedge.

Agent Continuity for Enterprises.

Context is handled – that was the entry ticket.
The next act is helping every fast-agentizing
company **actually put its agents to work**. The
platforms below keep their layer, and feed ours.
Echo becomes the layer their agents run on.



The floor is already priced near \$10B. The layer above it is where we're going.

Context continuity will let companies soar with agents. But what kind of society are we heading into?

Writing preserved **language**. Images preserved **faces**. Video let **the moment** exist apart from the room it happened in.

The human reach for permanence **kept widening** – and it defined the **society we live in**.

And today, Echo lets **the process by which thought forms** be preserved – and **transmitted, shared, connected, and remade**.

WHO IS BUILDING IT

Kobe Wang Founder & CEO

Former head of social product at **Roblox**, where social DAU passed **100M**; former PM at Renren, CMU. Three-time founder, **self-funding Echo**.



Carnegie Mellon University

6 ppl · 2 yr

Zero core-team attrition

Microsoft AI, Amazon, UC Berkeley, UCLA, BU, CMU

5 yr

Full-time all-in on Echo

Social → knowledge management → context management, end to end

100%

AI-native by default

Heavy Codex / Claude Code / OpenClaw users. No outsourcing

SF

Echo House · Marina, San Francisco

Half the team lives together

Agents carry us into a new era. We will build a new society.